

Mean Shift clustering algorithm

Lorenzo Agnolucci

Università degli Studi di Firenze
Dipartimento di Ingegneria dell'Informazione

Firenze, 27 Aprile 2020



- 1 Introduction
- 2 OpenMP
- 3 CUDA
 - Naive version
 - Tiling version
- 4 Experimental results
 - OpenMP
 - CUDA
 - Global comparison
- 5 Conclusions

- Mean Shift is a non-parametric clustering algorithm
- It is based on *Kernel Density estimation*
- The only parameter is the *bandwidth*
- $O(n^2)$ computational cost
- Common application in computer vision: image segmentation
- It is embarassingly parallel

- At each step a kernel function is applied to each point to make it shift towards the local maxima
- Most used kernel: Gaussian kernel

$$K(x) = e^{-\frac{x^2}{2\sigma^2}} \quad (1)$$

- New position x' where x has to be shifted is computed as:

$$x' = \frac{\sum_{x_i \in N(x)} K(\text{dist}(x, x_i)) x_i}{\sum_{x_i \in N(x)} K(\text{dist}(x, x_i))} \quad (2)$$

$N(x)$ is the neighborhood of x , a set of points for which $K(x_i) \neq 0$



- Algorithm stops when all points have stopped shifting, that is they have reached the local maxima

Algorithm 1 Mean shift core

```
function MEANSHIFT(originalPoints)  
  shiftedPoints  $\leftarrow$  originalPoints  
  while iterationIndex < MAX_ITERATIONS do  
    for each point p in shiftedPoints do  
      p  $\leftarrow$  SHIFTPoint(p, originalPoints)
```

Algorithm 2 Shift a single point

```
function SHIFTPoint(p, originalPoints)  
  shiftedP  $\leftarrow$  0  
  weight  $\leftarrow$  0  
  for each point x in originalPoints do  
    dist  $\leftarrow$  dist(p, x)  
    w  $\leftarrow$  GKernel(dist, BW)  
    shiftedP  $\leftarrow$  shiftedP + w * x  
    weight  $\leftarrow$  weight + w  
  return shiftedP / weight
```



Algorithm 3 OpenMP Mean shift core

```
function OPENMPMEANSHIFT(originalPoints)  
  shiftedPoints  $\leftarrow$  originalPoints  
  while iterationIndex < MAX_ITERATIONS do  
    #pragma parallel for schedule(static)  
    for each point p in shiftedPoints do  
      p  $\leftarrow$  SHIFTPPOINT(p, originalPoints)
```

- just a pragma directive
- static scheduling: loop divided statically in chunks of equal size

Algorithm 4 Shift a single point

```
function SHIFTPPOINT(p, originalPoints)  
  shiftedP  $\leftarrow$  0  
  weight  $\leftarrow$  0  
  for each point x in originalPoints do  
    dist  $\leftarrow$  dist(p, x)  
    w  $\leftarrow$  GKernal(dist, BW)  
    shiftedP  $\leftarrow$  shiftedP + w * x  
    weight  $\leftarrow$  weight + w  
  return shiftedP / weight
```

Taking advantage of GPUs:

- One thread for each point to shift
- Coalescing of accesses to memory:
 - ◆ Points stored as a *Structure of Arrays*

$$[x_1, \dots, x_n, y_1, \dots, y_n, z_1, \dots, z_n] \quad (3)$$

- ◆ Access to the array in the form of:

$$blockDim.x * blockIdx.x + threadIdx.x \quad (4)$$

Algorithm 5 CUDA Naive version Mean Shift core

```
function NAIVECUDAMS(originalPoints)  
  shiftedPoints  $\leftarrow$  originalPoints  
  while iterationIndex < MAX_ITERATIONS do  
    NAIVEKERNEL(shiftedPoints, originalPoints)
```

Algorithm 6 CUDA Naive version Kernel

```
function NAIVEKERNEL(shiftedPts, originalPts)  
  tx  $\leftarrow$  threadIdx.x  
  bx  $\leftarrow$  blockIdx.x  
  idx  $\leftarrow$  bx * blockDim.x + tx  
  if idx < |originalPts| then  
    shiftedP  $\leftarrow$  0  
    weight  $\leftarrow$  0  
    p  $\leftarrow$  shiftedPts[idx]  
    for each point x in originalPts do  
      dist  $\leftarrow$  dist(p, x)  
      w  $\leftarrow$  GKernal(dist, BW)  
      shiftedP  $\leftarrow$  shiftedP + w * x  
      weight  $\leftarrow$  weight + w  
    shiftedPts[idx]  $\leftarrow$  shiftedP / weight
```

A further optimization is possible:

- Each thread reads $O(n)$ points from global memory to compute the shift
- Shared Memory can be exploited with the Tiling pattern
- $TILE_WIDTH = BLOCK_DIM$
- $O(n)$ accesses reduced to $O(n/TILE_WIDTH)$

Algorithm 7 CUDA Tiling version Mean Shift core

```
function TILINGCUDAMS(originalPoints)  
    shiftedPoints  $\leftarrow$  originalPoints  
    while iterationIndex < MAX_ITERATIONS do  
        TILINGKERNEL(shiftedPoints, originalPoints)
```

Algorithm 8 CUDA Tiling version Kernel

```
function TILINGKERNEL(shiftedPts, originalPts)  
    tx  $\leftarrow$  threadIdx.x  
    bx  $\leftarrow$  blockIdx.x  
    idx  $\leftarrow$  bx * blockDim.x + tx  
    tile  $\leftarrow$  SharedMemArray[TILE_WIDTH]  
    shiftedP  $\leftarrow$  0  
    weight  $\leftarrow$  0  
    for tileIter < numTiles do  
        tileIdx  $\leftarrow$  tileIter * TILE_WIDTH + tx  
        if tileIdx < |originalPts| then  
            tile[tx]  $\leftarrow$  originalPts[tileIdx]  
        else  
            tile[tx]  $\leftarrow$  nullPoint  
        __syncthreads()  
        if idx < |originalPts| then  
            p  $\leftarrow$  shiftedPts[idx]  
            for i with i < TILE_WIDTH do  
                x  $\leftarrow$  tile[i]  
                if x! = nullPoint then  
                    dist  $\leftarrow$  dist(p, x)  
                    w  $\leftarrow$  GKernel(dist, BW)  
                    shiftedP  $\leftarrow$  shiftedP + w * x  
                    weight  $\leftarrow$  weight + w  
            __syncthreads()  
        if idx < |originalPts| then  
            shiftedPts[idx]  $\leftarrow$  shiftedP / weight
```

▷ End of loading

▷ End of computing

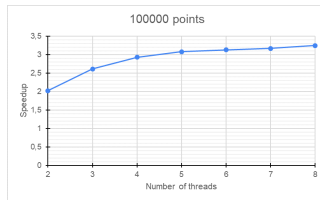
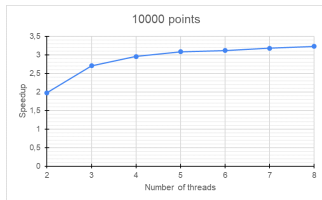
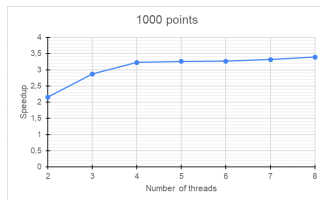
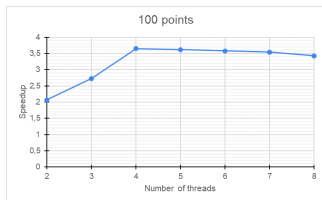
- Performances compared with the speedup metric, computed as:

$$S = \frac{t_S}{t_P} \quad (5)$$

- Tests executed on a machine with:
 - ◆ OS: Ubuntu 18.04 LTS
 - ◆ CPU: Intel Core i7-8565U 1.8GHz up to 4.6GHz with Turbo Boost, 4 cores/8 threads
 - ◆ RAM: 16 GB DDR4
 - ◆ GPU: NVidia GeForce MX250 2GB with CUDA 10.1
- Each time is the average of 5 experiments for the sequential and OpenMP versions and of 15 experiments for both CUDA implementations

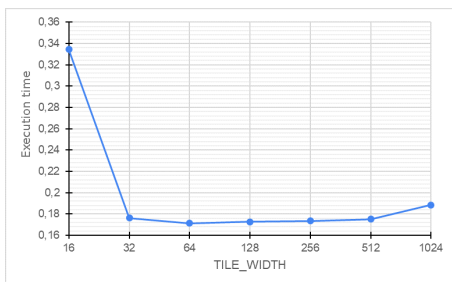
- The implementations have been evaluated on gaussian distributions composed by respectively 100, 1000, 10000, 100000 and 250000 3D points
- bandwidth set to 2
- *MAX_ITERATIONS* constant set to 10, which has been empirically estimated to be enough to make all points converge to the local maxima

Number of threads gradually increased for each dataset



Not tested on 250k points dataset: too long execution time

Execution time for 10000 points for the CUDA tiling implementation varying *TILE_WIDTH*



| TILE_WIDTH | CUDA Tiling |
|------------|-------------------|
| 16 | 0.334405 s |
| 32 | 0.176171 s |
| 64 | 0.171354 s |
| 128 | 0.172666 s |
| 256 | 0.173306 s |
| 512 | 0.175091 s |
| 1024 | 0.18845 s |

| Dim | CUDA Naive | CUDA Tiling | Speedup |
|------------|-------------------|--------------------|----------------|
| 100 | 0.000401 s | 0.000437 s | 1.09 |
| 1000 | 0.003044 s | 0.003359 s | 1.11 |
| 10000 | 0.171354 s | 0.191216 s | 1.12 |
| 100000 | 15.79 s | 18.29 s | 1.16 |
| 250000 | 100.38 s | 121.10 s | 1.20 |

| Dim | Sequential | CUDA Tiling | Speedup |
|--------|------------|-------------|----------|
| 100 | 0.005384 s | 0.000401 s | 13.43 |
| 1000 | 0.475485 s | 0.003044 s | 156.21 |
| 10000 | 49.29 s | 0.171354 s | 287.64 |
| 100000 | 4560.37 s | 15.79 s | 288.80 |
| 250000 | † 27768 s | 100.38 s | † 276.61 |

- † time has been estimated with a quadratic regression (due to the $O(n^2)$ computational cost)

Global comparison between sequential, OpenMP and Tiling CUDA best results:

- Greatest speedups with CUDA, at the expense of a more complicated implementation
- OpenMP lets to reach noticeable speedups with a simple implementation (just a directive)

| Dim | Sequential | OpenMP | OpenMP Speedup | CUDA Tiling | CUDA Speedup |
|--------|------------|------------|----------------|-------------|--------------|
| 100 | 0.005384 s | 0.001473 s | 3.66 | 0.000401 s | 13.43 |
| 1000 | 0.475485 s | 0.140161 s | 3.39 | 0.003044 s | 156.21 |
| 10000 | 49.29 s | 15.25 s | 3.24 | 0.171354 s | 287.64 |
| 100000 | 4560.37 s | 1403.78 s | 3.25 | 15.79 s | 288.80 |
| 250000 | † 27768 s | † 8720 s | † 3.18 | 100.38 s | † 276.61 |

Table: † times have been estimated with a quadratic regression

- The embarrassingly parallel structure of Mean Shift makes it suitable for parallel implementations
- OpenMP has an excellent speedup and development cost ratio
- CUDA makes Mean Shift applicable to datasets intractable with a CPU